



MSML610: Advanced Machine Learning

Class mechanics

Instructor: GP Saggese, PhD - gsaggese@umd.edu

MSML610

- **MSML610**
- Class map

Books of the class

- The goal is to make the slides self-sufficient
- For each lesson, I recommend a few books I liked about that topic
 - Simple -> Medium -> Hardcore
- Simple
 - Burkov: "Machine Learning Engineering" (2020)
 - Burkov: "The Hundred-Page Machine Learning Book" (2019)
- Medium
 - Abu-Mostafa et al.: "Learning From Data" (2012)
 - Martin: "Bayesian Analysis with Python" (2nd ed, 2021)
 - Russell et al.: "Artificial Intelligence: A Modern Approach" (4th ed, 2020)
- Hardcore
 - Hastie et al.: "The Elements of Statistical Learning" (2nd ed, 2009)
 - Koller et al.: "Probabilistic Graphical Models: Principles and Techniques" (2009)
 - Murphy: Machine Learning: "A Probabilistic Perspective" (2012)
 - Sutton et al.: "Reinforcement Learning: An Introduction" (2nd ed, 2018)

Invariants of a class lecture

- Focus on intuition over math (unless necessary)
- Emphasize realistic assumptions and numerical methods
 - Analytical solutions are so 1800s
- Each lesson includes an interactive Jupyter notebook tutorial to foster intuition
 - Tutorials are mainly done at home
 - Videos of each tutorial will be added over time
- Lessons alternate between slides, whiteboard, and tutorials
- 2:45 hours of slides
 - 50 slides
 - 10 break
 - 50 slides
 - 10 break
 - 45 slides (Topic refresher!)

Grading

- Class Participation (10%): Contributions to discussions and engagement
- Quizzes (40%): 4-5 quizzes to force people to study as we go
 - Multi-choice quizzes
 - To make sure everyone studies during the semester and doesn't cram
- Final Project (50%): A comprehensive application of course concepts

Class projects

- There is a list of topics you can pick from, e.g.,
 - LLMs
 - Deep learning
 - Big data
 - Statistical learning
 - ...
- Different levels of difficulty
- Each project is individual
- Content of each project XYZ
 - Write a blog entry about XYZ
 - Create a Jupyter notebook about XYZ
 - Implement a meaningful use case about using XYZ
 - Commit code to GitHub and contribute to open-source repo

Links

- Syllabus
- FAQs
- Project specs

- Announcements

Class map

- MSML610
- **Class map**

1. Intro

- A map of machine learning
- What is Artificial Intelligence
 - AI
 - ML
 - AI vs ML vs Deep-learning
 - The foundation of AI
 - Brief history of AI
 - AI state of the art
 - Risks and benefits of AI

2. Techniques

- Paradigms
- Techniques
 - Machine learning in practice
 - How to do research
 - Simple is better
 - Research methodology
 - Pipeline organization
 - Input processing
 - Learning algorithms
 - Gradient descent
 - Stochastic gradient descent
 - Performance metrics
 - Precision and recall
 - Model selection
 - Aggregation
 - Bagging
 - Boosting
 - Stacking

3. Knowledge Representation

- Knowledge Representation
 - Basics of Knowledge Representation
 - Examples of Logic
 - Logical Agents
 - Ontologies
 - Reasoning in Ontologies
- Propositional logic
- First-order Logic
- Non-classical Logics
- Description Logics
 - Semantic Web
- Advanced topics

4. Machine Learning Models

- Models
 - Naive Bayes
 - Decision trees
 - Random forests
 - Linear models
 - Perceptron
 - Logistic regression
 - LDA, QDA
 - Kernel methods
 - Support vector machines
 - Similarity-based models
 - Clustering
 - Anomaly detection

5. Machine Learning Theories

- Is learning possible?
 - Training vs Testing
 - Growth function
 - The VC dimension
 - Regularization
 - Bias vs variance
 - Learning curves
 - Learn-validation approach

6. Bayesian statistics

- Quantifying uncertainty
- Probabilistic reasoning

7. Probabilistic programming

- Probabilistic programming
 - Probability theory
 - Single-parameter inference
 - How to choose priors
 - Communicating a Bayesian analysis
 - Probabilistic programming
 - Posterior-based decisions
 - Gaussians all the way down
 - Posterior predictive checks
 - Robust inference
 - Groups comparison
 - Hierarchical models
 - Simple linear model
 - Variable variance
 - Hierarchical linear regression
 - Multiple linear regression
 - Comparing models
 - Posterior predictive checks
 - The balance between simplicity and accuracy
 - Measures of predictive accuracy
 - Information criteria
 - Cross-validation

8. Reasoning over time

- Reasoning over time
- HMMs
- Markov random fields
- Markov logic network
- State space models and Kalman filter
 - g-h filter
 - Discrete Bayes filter
- Dynamic Bayesian networks
- State space model
- Variational Inference
 - Expectation-Maximization (EM) Algorithm

9. Causal inference

- Causal AI
 - Why Causal AI?
 - Concepts in Causal AI
 - Variables
 - Paths
 - The Ladder of Causation
 - Correlation vs causation models
- Business processes around data modeling
 - Modeling processes
 - Roles

10. Timeseries forecasting

- Time Series
 - Basic definition
 - Time series operators
 - Time series decomposition
- Classical Methods
 - Simple models for stochastic process
 - Autoregressive models
 - Moving average models
 - ARMA(p , q) process
 - ARIMA model
 - ARCH model

11. Probabilistic deep learning

- Neural networks
 - Biological inspiration
 - Neural networks
- Advanced Neural Network Architectures
 - Convolutional networks
 - Recurrent Neural Networks (RNNs)
 - Deep learning learning algorithms
 - Deep learning architectures
- Fundamentals of Deep Learning
- Training Deep Neural Networks
- Interpretability and Explainability
- Deep Generative Models
- Bayesian Deep Learning
- Deep Probabilistic Models
- Uncertainty Quantification
- Probabilistic Programming and Inference
- Modern Research Frontiers
- Bonus Topics

12. Reinforcement learning

- Sequential decision problems
 - Utilities over time
 - Algorithms for MDPs
- Reinforcement learning
 - Passive reinforcement learning
 - Active reinforcement learning
 - Generalization in reinforcement learning
 - Policy search
- Fundamentals
- Classical Methods
- Exploration Strategies
- Policy Gradient Methods
- Value Function Approximation
- Deep Reinforcement Learning
- Model-Based Reinforcement Learning
- Advanced Topics
- Applications

Refresher: Probability

- Probability
 - Probability definition
 - Probability measure
 - Independent events
 - Conditional probability
 - Law of total probability
 - Bayes theorem
- Random variables
 - Random variables
 - CDF, PMF, PDF of Random Variables
 - Joint distributions
 - Marginal distributions
 - Independent RVs
 - Conditional PDF RVs
- Mathematical expectation of RVs
 - Mean
 - Variance and covariance
 - Statistics of RVs
- Probability inequalities
- Statistical Inference
 - Definitions

Refresher probability distributions

- Interesting RVs
 - Bernoulli
 - Binomial
 - Gaussian
 - Log-Normal
 - Poisson
 - Chi-square
 - Student's t-distribution
- Probability inequalities

Refresher linear algebra

- Linear algebra
 - Vector and vector spaces
 - Affine spaces
 - Vectors and matrices
 - Linear functions
 - Connections between Machine Learning and Linear Algebra

Refresher information theory

- Information theory
 - Entropy
 - Kullback-Leibler divergence
 - Connections between Information Theory and ML

Refresher game theory

- Game theory
 - Connections between Machine Learning and Game Theory

Refresher: numerical optimization

- Optimization / numerical methods

Refresher: stochastic processes

- Stochastic processes